Q6 Exercise 2.9 :

Solution: UCB action selection is given by:

$$A_t = \underset{a}{\operatorname{argmax}} \left[Q_t(a) + c \sqrt{\frac{\ln t}{N_t(a)}} \right]$$

Here, c is greater than 0 and controls the degree of exploration.

Initially, $N_t(a) = 0 \quad \forall a$

⇒ for any action a, until it has been selected once $N_t(a)$ remains 0.

⇒ All the actions are explored one by one in initial 10 time steps (in our case), no matter whatever the value of c is.

Once it has explored all the actions, it chooses the greedy one hence average reward over 2000 runs increases.

→ One can view the square-root quantity as the variance or the uncertainty in the estimate of a's value as stated in Sutton.

⇒ More is the square root quantity, more the uncertainty in the estimate of an action.

After 11th step, $N(a)$ of the action chosen at 10th step increases and uncertainty decreases. But at the same time, uncertainty for other actions increases as $\ln t$ increases. Therefore, UCB algorithm continues to explore at subsequent step, hence average reward decreases.

Page No

Scanned by CamScanner

Now, for different values of c, one may see the different graph.
This is b/c as c increases, ~~uncertainty~~ the ~~c~~ actions with higher uncertainty gets ~~pulled~~ selected.
Therefore, if c=1, the spike is less prominent as it chooses an action which has a good estimate rather than high uncertainty.

Similarly, when c=4, actions ~~which~~ which are not greedy gets selected after 11th step b/c they have higher uncertainty